

GenoPHI

Genotype-to-Phenotype Phage-Host Interaction Prediction

GenoPHI is a Python package for machine learning-based prediction of genotype-phenotype relationships using whole-genome sequence data. Originally designed for phage-host interaction prediction, GenoPHI supports both binary interaction prediction and regression tasks for any microbial phenotype. The package implements protein family-based and k -mer-based approaches to extract genomic features from amino acid sequences and predict phenotypes using CatBoost gradient boosting models.

License MIT python 3.8+

Workflow Overview

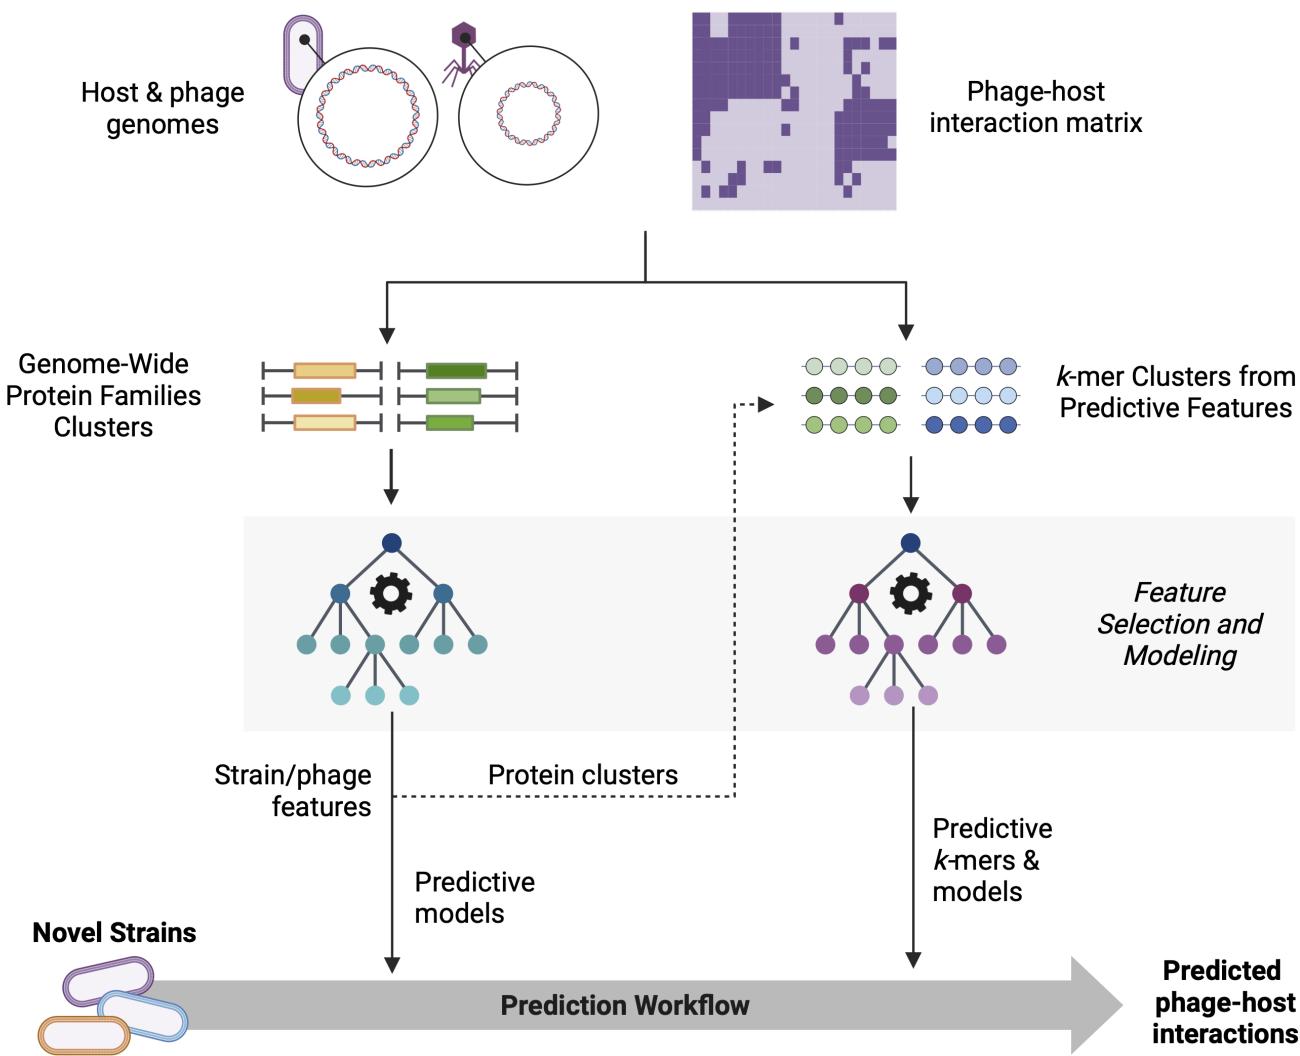


Figure 1: GenoPHI workflow schematic showing the main analysis pipelines: Protein family-based workflow, K-mer-based workflow, and Predictive protein k-mer workflow. Each pathway includes feature extraction, selection, model training, and prediction steps.

Table of Contents

- [Features](#)
- [Installation](#)
- [Quick Start](#)
- [Usage](#)
 - [CLI Commands Overview](#)
 - [Workflows](#)
- [Input Data Formats](#)
- [Feature Selection Methods](#)
- [Performance Metrics](#)
- [Output Directory Structure](#)
- [Python API](#)
- [Advanced Usage](#)
- [Troubleshooting](#)
- [Best Practices](#)
- [Citation](#)
- [Contributing](#)
- [License](#)
- [Support](#)
- [Acknowledgments](#)

Features

Protein Family-Based Analysis

1. **MMSeqs2 Clustering:** Cluster protein sequences into protein families based on sequence similarity
2. **Feature Table Generation:** Create presence-absence matrices of protein families across genomes and consolidate into predictive features based on co-occurrence across genomes
3. **Feature Selection:** Identify predictive protein families (multiple available methods: RFE, SHAP, SHAP-RFE, ANOVA, Chi-squared, Lasso)
4. **Model Training:** Train CatBoost models with hyperparameter optimization
5. **Phenotype Prediction:** Predict interactions, resistance, or other phenotypes for new genomes
6. **Feature Annotation:** Identify predictive protein sequences from predictive features

K-mer-Based Analysis

1. **K-mer Feature Extraction:** Generate k -mer features from protein sequences with or without gene family context
2. **Predictive K-mer Workflow:** Extract k -mers specifically from predictive protein families identified in protein family analysis
3. **Feature Selection & Modeling:** Apply same robust feature selection and modeling pipelines
4. **Flexible K-mer Lengths:** Support for single k -mer length or ranges (e.g., 3-6)

Application Modes

- **Phage-Host Interaction Prediction:** Binary prediction of infection outcomes between phages and bacterial strains
- **Single-Strain Phenotype Prediction:** Predict strain-level phenotypes (e.g., antibiotic resistance, growth rate) without requiring phage data

- **Regression Tasks:** Predict continuous phenotypes (e.g., infection efficiency, metabolic rates)
- **General Feature-Based Modeling:** Use any feature table with a phenotype column for custom applications

Advanced Capabilities

- **Dynamic Feature Weighting:** Account for feature frequency distributions to handle imbalanced features
- **Clustering-Based Selection:** Use HDBSCAN or hierarchical clustering for intelligent feature grouping
- **Multiple Feature Selection Methods:** RFE, SHAP-RFE, SelectKBest, Chi-squared, Lasso, SHAP
- **Comprehensive Performance Metrics:** AUC-ROC, Precision-Recall, MCC, F1-score, Accuracy
- **SHAP Interpretability:** Feature importance analysis and visualization for model explainability
- **Bootstrapping Support:** Robust model evaluation with multiple train-test splits

Installation

Prerequisites

External Dependency: GenoPHI requires MMseqs2 for protein sequence clustering and assignment.

Install via conda/mamba:

```
conda install -c bioconda mmseqs2
# or
mamba install -c bioconda mmseqs2
```

For other installation methods, see the [MMSeqs2 Wiki](#).

System Requirements

Minimum Requirements:

- Python 3.8 or higher
- 8 GB RAM
- 4 CPU cores
- 10 GB free disk space

Recommended for Large Datasets:

- Python 3.10+
- 32+ GB RAM
- 8+ CPU cores
- 50+ GB free disk space (depending on dataset size)

Tested Operating Systems:

- Linux (Ubuntu 20.04+, CentOS 7+)

Virtual Environment (Recommended)

Create and activate a conda environment:

```
conda create -n genophi python=3.10
conda activate genophi
```

Install GenoPHI

Clone and install from GitHub:

```
git clone https://github.com/Noonanav/GenoPHI.git
cd GenoPHI
pip install -e .
```

For development with optional dependencies:

```
pip install -e ".[dev]"
```

Verify Installation

Test that GenoPHI is properly installed:

```
# Check GenoPHI version
genophi --version

# Verify MMseqs2 is accessible
mmseqs version

# Run basic help command
genophi --help
```

Quick Start

GenoPHI provides a unified command-line interface accessible through the `genophi` command:

```
# View available commands
genophi --help

# Get help for a specific command
genophi protein-family-workflow --help
```

Recommended Default Run

For most phage-host interaction prediction tasks, use these recommended settings:

```
genophi protein-family-workflow \
--input_strain strain_fastas/ \
--input_phage phage_fastas/ \
--phenotype_matrix interactions.csv \
--output results/ \
--threads 8 \
--num_features 100 \
--num_runs_fs 25 \
--num_runs_modeling 50 \
--method rfe \
--use_clustering \
--cluster_method hierarchical \
--n_clusters 20 \
--filter_type strain
```

Key Parameters Explained:

- `--num_features 100`: Select top 100 features (adjust based on dataset size)
- `--num_runs_fs 25`: 25 iterations for robust feature selection
- `--num_runs_modeling 50`: 50 modeling runs for reliable performance estimates
- `--method rfe`: Recursive Feature Elimination (balanced performance)
- `--use_clustering`: Enable hierarchical clustering of features
- `--filter_type strain`: Split data by strain to test generalization

For single-strain phenotypes (no phage data):

```
genophi protein-family-workflow \
--input_strain strain_fastas/ \
--phenotype_matrix phenotypes.csv \
--output results/ \
--threads 8 \
--sample_column strain \
--phenotype_column resistance \
--filter_type none
```

Usage

CLI Commands Overview

GenoPHI provides the following main commands:

Command	Description
<code>protein-family-workflow</code>	Recommended basic workflow: Complete protein family-based workflow
<code>full-workflow</code>	Protein families → k -mers from predictive proteins

Command	Description
kmer-workflow	Complete k -mer-based workflow from all proteins
cluster	Generate protein family clusters and feature tables
select-features	Perform feature selection on any feature table
train	Train predictive models on selected features
predict	Predict phenotypes using trained models
select-and-train	Feature selection + modeling from any feature table
assign-features	Assign features to new genomes
assign-predict	Assign features and predict (protein families)
annotate	Annotate predictive features with functional info
kmer-assign-features	Assign k -mer features to new genomes
kmer-assign-predict	Assign k -mer features and predict

Workflows

1. Protein Family Workflow (Recommended)

The primary workflow for most applications. Performs complete protein family clustering, feature selection, and modeling.

Complete Workflow

```
genophi protein-family-workflow \
--input_strain strain_fastas/ \
--input_phage phage_fastas/ \
--phenotype_matrix interactions.csv \
--output results/ \
--threads 8 \
--num_features 100 \
--num_runs_fs 25 \
--num_runs_modeling 50 \
--method rfe \
--filter_type strain
```

Output Structure:

```
results/
└── clustering/          # MMseqs2 databases and clusters
└── feature_tables/      # Generated feature tables
└── feature_selection/    # Selected features and occurrence counts
└── modeling_results/     # Models and performance metrics
```

```
|   └── cutoff_5/, cutoff_10/, ...
|   └── model_performance/    # Summary plots
|       └── models/          # Trained models for prediction
|   └── workflow_report.txt  # Performance and timing summary
```

Single-Strain Phenotype Prediction

For strain-level phenotypes (no phage data required):

```
genophi protein-family-workflow \
--input_strain strain_fastas/ \
--phenotype_matrix strain_phenotypes.csv \
--output results/ \
--threads 8 \
--sample_column strain \
--phenotype_column antibiotic_resistance \
--task_type classification \
--filter_type none
```

Phenotype Matrix Format:

```
strain,antibiotic_resistance
Strain_001,1
Strain_002,0
Strain_003,1
```

For regression:

```
--task_type regression \
--phenotype_column growth_rate
```

2. Full Workflow: Protein Families → Predictive *K*-mers

This workflow first identifies predictive protein families, then extracts *k*-mers specifically from those families for refined modeling. This combines the interpretability of protein families with the resolution of *k*-mer analysis.

```
genophi full-workflow \
--input_strain strain_fastas/ \
--input_phage phage_fastas/ \
--phenotype_matrix interactions.csv \
--output results/ \
--k 5 \
--threads 8
```

Workflow Steps:

1. Cluster proteins into families
2. Perform feature selection on protein families
3. Extract k -mers from predictive protein families only
4. Train models on k -mer features
5. Generate annotations for predictive k -mers

3. K-mer Workflow (*De Novo*)

Generate k -mer features from all proteins without prior protein family analysis:

```
genophi kmer-workflow \  
  --input_strain strain_fastas/ \  
  --input_phage phage_fastas/ \  
  --phenotype_matrix interactions.csv \  
  --output kmer_results/ \  
  --k 4 \  
  --threads 8 \  
  --num_features 100 \  
  --num_runs_fs 25 \  
  --num_runs_modeling 50
```

K-mer Specific Parameters:

- `--k 4`: K -mer length (default: 4)
- `--k_range`: Generate k -mers from length 3 to k
- `--one_gene`: Include features with only one gene (default: False)

Advanced Options:

```
--use_dynamic_weights \      # Apply dynamic weighting  
--weights_method inverse_frequency \    # Weighting method  
--use_feature_clustering \    # Pre-filter by cluster presence  
--feature_n_clusters 20      # Number of feature clusters
```

4. Modular Step-by-Step Workflows

Step 1: Clustering and Feature Generation

```
genophi cluster \  
  --input_strain strain_fastas/ \  
  --input_phage phage_fastas/ \  
  --phenotype_matrix interactions.csv \  
  --output clustering_results/ \  
  --min_seq_id 0.4 \
```

```
--coverage 0.8 \
--sensitivity 7.5 \
--threads 8
```

Clustering Parameters:

- `--min_seq_id 0.4`: Minimum sequence identity (range: 0-1)
- `--coverage 0.8`: Minimum coverage (range: 0-1)
- `--sensitivity 7.5`: MMseqs2 sensitivity (higher = more sensitive, slower)

Step 2: Feature Selection

Feature selection works on **any** feature table with a phenotype column:

```
genophi select-features \
--input feature_table.csv \
--output feature_selection/ \
--method rfe \
--num_features 100 \
--num_runs 25 \
--filter_type strain \
--phenotype_column interaction \
--threads 8
```

Feature Selection Methods:

- `rfe`: Recursive Feature Elimination (recommended)
- `shap_rfe`: RFE using SHAP values
- `select_k_best`: ANOVA F-test (fast)
- `chi_squared`: Chi-squared test
- `lasso`: L1 regularization
- `shap`: Direct SHAP importance

Advanced Selection Options:

```
--use_dynamic_weights \           # Handle imbalanced features
--weights_method inverse_frequency \ # Weighting strategy
--use_clustering \                 # Group correlated features
--cluster_method hierarchical \    # Clustering algorithm
--n_clusters 20                  # Number of clusters
```

Step 3: Model Training

Train models from selected features or any feature table:

```
genophi train \  
  --input_dir feature_selection/filtered_feature_tables \  
  --output models/ \  
  --num_runs 50 \  
  --phenotype_column interaction \  
  --threads 8
```

For regression tasks:

```
--task_type regression \  
--phenotype_column efficiency
```

Step 4: Select-and-Train (Combined)

Run feature selection and modeling together from **any feature table**:

```
genophi select-and-train \  
  --input custom_feature_table.csv \  
  --output results/ \  
  --method rfe \  
  --num_features 100 \  
  --num_runs_fs 25 \  
  --num_runs_modeling 50 \  
  --phenotype_column your_phenotype \  
  --sample_column your_sample_id \  
  --threads 8
```

This command is flexible and works with:

- Protein family features
- *K*-mer features
- Custom features
- Any feature table with any phenotype / output column

5. Prediction Workflows

Assign Features and Predict (Protein Families)

```
genophi assign-predict \  
  --input_dir new_strains/ \  
  --mmseqs_db results/tmp/strain/mmseqs_db \  
  --clusters_tsv results/strain/clusters.tsv \  
  --feature_map results/strain/features/selected_features.csv \  
  --model_dir results/modeling_results/cutoff_* \  
  --phage_feature_table results/phage/features/feature_table.csv \  
  --output_dir results/predictions/
```

```
--output predictions/ \
--genome_type strain
```

For new phages:

```
--input_dir new_phages/ \
--mmseqs_db results/tmp/phage/mmseqs_db \
--clusters_tsv results/phage/features/selected_features.csv \
--strain_feature_table results/strain/features/feature_table.csv \
--genome_type phage
```

K-mer Assignment and Prediction

```
genophi kmer-assign-predict \
--input_dir new_strains/ \
--mmseqs_db results/tmp/strain/mmseqs_db \
--clusters_tsv results/strain/clusters.tsv \
--feature_map results/strain/features/selected_features.csv \
--filtered_kmers kmer_results/kmer_tables/filtered_kmers.csv \
--aa_sequence_file kmer_results/kmer_tables/aa_sequences.faa \
--model_dir kmer_results/modeling/modeling_results/cutoff_* \
--output predictions/ \
--genome_type strain
```

6. Feature Annotation

Identifies proteins associated predictive protein families or *k*-mers and merges with functional information:

```
genophi annotate \
--selected_features
feature_selection/filtered_feature_tables/select_feature_table_cutoff_3.csv \
--feature_map strain/features/feature_map.csv \
--clusters_tsv clustering_results/clustering/selected_features.csv \
--annotation_table annotations.csv \
--aa_sequence_file all_sequences.faa \
--output annotations/ \
--feature_type strain
```

Input Data Formats

FASTA Files

Protein sequences in FASTA format (.faa files):

```
>protein_id_1  
MKTAYIAKQRQISFVKSHFSRQLEERLGLIEQAPILSRVGDGTQDNLSGAEKAVQVKVKALPDAQFEVHSLAKWKRQ...  
>protein_id_2  
MRISTTITTTITTTGNGAG...
```

Important: Protein IDs must be unique across all genomes. If duplicates exist, GenoPHI will automatically prefix them with genome names.

Phenotype Matrix

Phage-Host Interactions

Binary classification (infection/no infection):

```
strain,phage,interaction  
Strain_001,Phage_A,1  
Strain_001,Phage_B,0  
Strain_002,Phage_A,1
```

Regression (infection efficiency):

```
strain,phage,efficiency  
Strain_001,Phage_A,0.85  
Strain_001,Phage_B,0.12  
Strain_002,Phage_A,0.93
```

Single-Strain Phenotypes

Classification:

```
strain,antibiotic_resistance  
Strain_001,1  
Strain_002,0  
Strain_003,1
```

Regression:

```
strain,growth_rate  
Strain_001,0.42  
Strain_002,0.38  
Strain_003,0.51
```

Column Names: Use `--strain_column`, `--phage_column`, `--sample_column`, and `--phenotype_column` to specify your column names.

Feature Selection Methods

Method	Description	Best For	Speed
RFE (recommended)	Recursive Feature Elimination	General use, balanced performance	Medium
SHAP-RFE	RFE using SHAP values	Model-agnostic importance	Slow (High RAM)
SelectKBest	ANOVA F-test	Fast screening, linear relationships	Fast
Chi-squared	χ^2 test for independence	Categorical features	Fast
Lasso	L1 regularized regression	Sparse models, multicollinearity	Fast
SHAP	Shapley Additive Explanations	Direct feature importance	Slow (High RAM)

Dynamic Weighting

Handle imbalanced feature distributions:

```
--use_dynamic_weights \
--weights_method inverse_frequency # or log10, balanced
```

When to use:

- Features with highly variable occurrence frequencies
- Some features present in most genomes, others very rare
- Imbalanced positive/negative examples

Clustering-Based Selection

Group correlated features for more robust selection:

```
--use_clustering \
--cluster_method hierarchical \ # or hdbscan
--n_clusters 20
```

HDBSCAN Options:

```
--cluster_method hdbscan \
--min_cluster_size 5 \
```

```
--min_samples 5 \
--cluster_selection_epsilon 0.0
```

Performance Metrics

Classification Metrics

- **AUC-ROC (Area Under ROC Curve):** Overall discriminative ability $\text{AUC} = \int_0^1 \text{TPR}(\text{FPR}) \, d\text{FPR}$
- **MCC (Matthews Correlation Coefficient):** Balanced metric for all confusion matrix elements $\text{MCC} = \frac{\text{TP} \times \text{TN} - \text{FP} \times \text{FN}}{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}}$
- **Precision:** Proportion of true positives among predicted positives $\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall (Sensitivity):** Proportion of true positives among actual positives $\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1 Score:** Harmonic mean of precision and recall $\text{F1} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$
- **Accuracy:** Overall correct predictions $\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$

Regression Metrics

- **RMSE:** Root Mean Squared Error
- **MAE:** Mean Absolute Error
- **R²:** Coefficient of determination

Visualization Outputs

GenoPHI generates comprehensive visualizations:

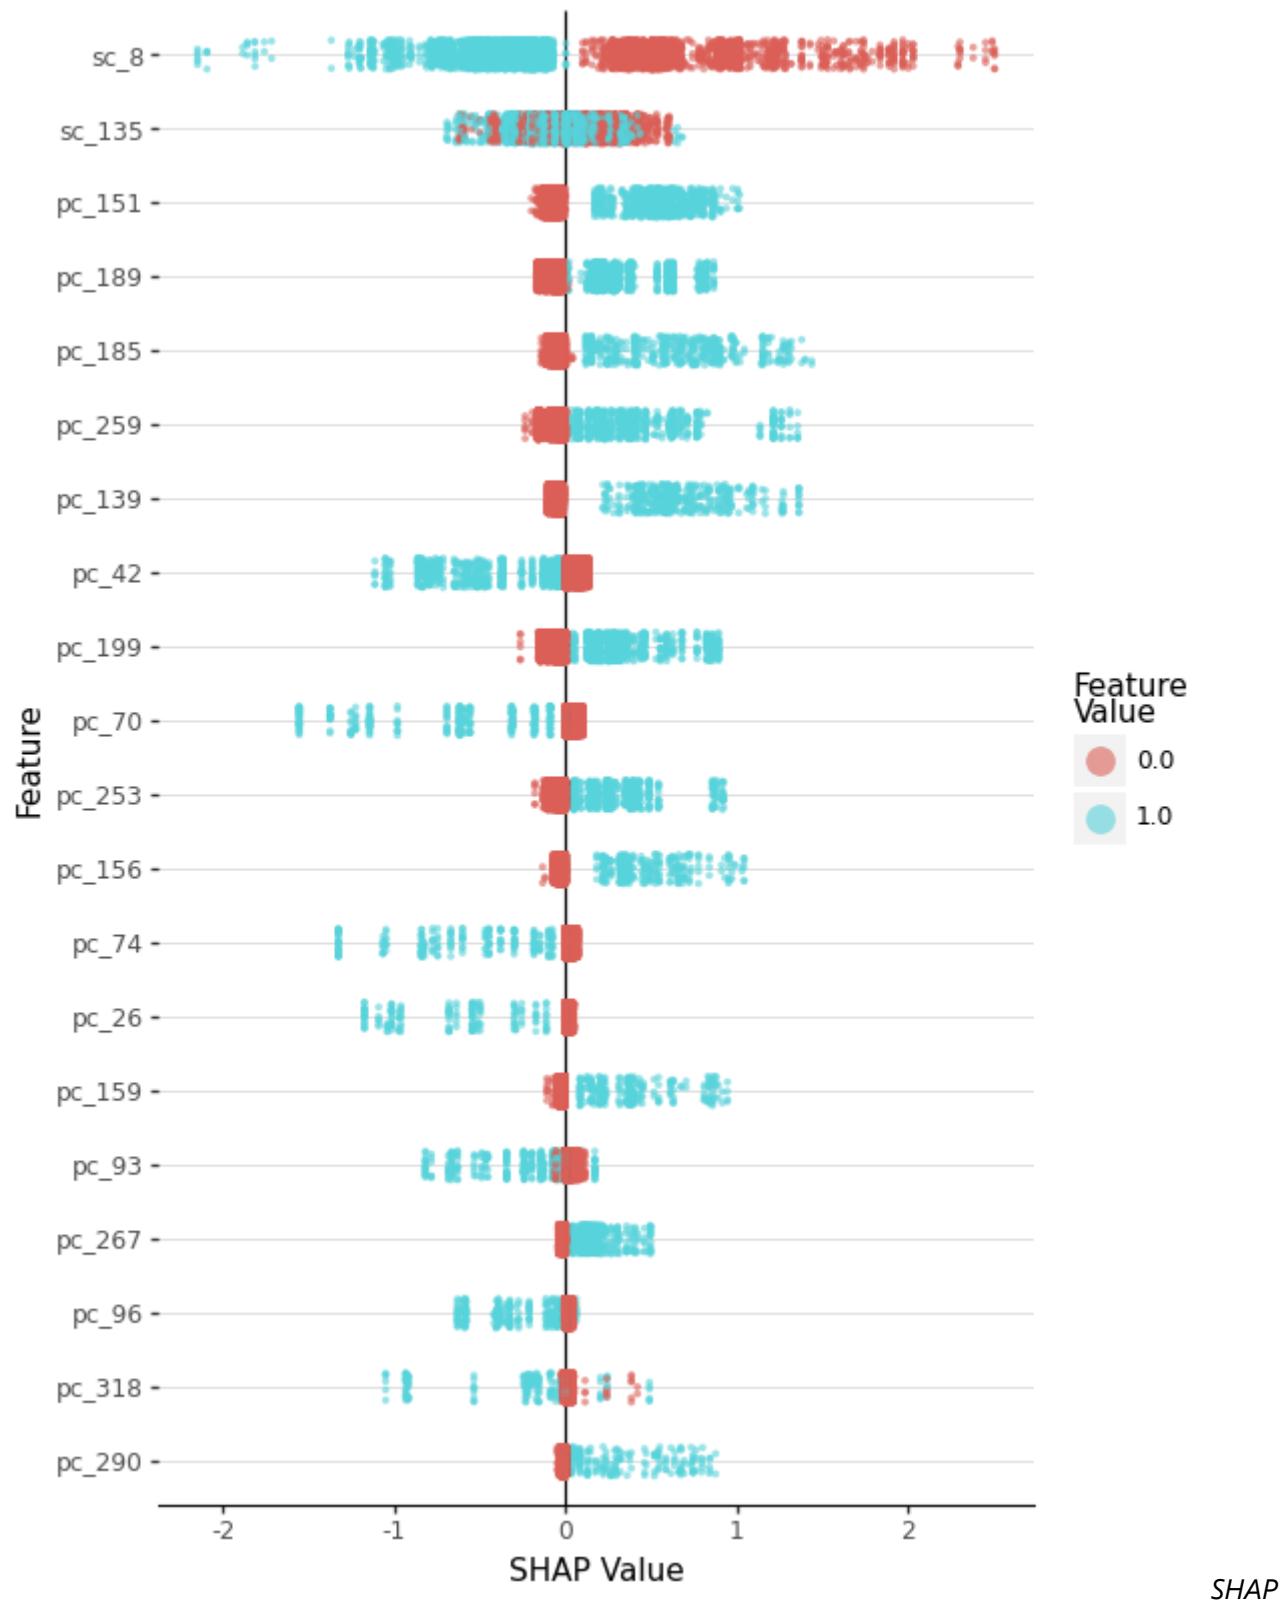
Per-Run Plots (`modeling_results/cutoff_*/run_*/`):

- Confusion matrices (classification)
- ROC curves with AUC scores
- Precision-Recall curves
- SHAP feature importance bar plots
- SHAP value scatter plots (beeswarm)

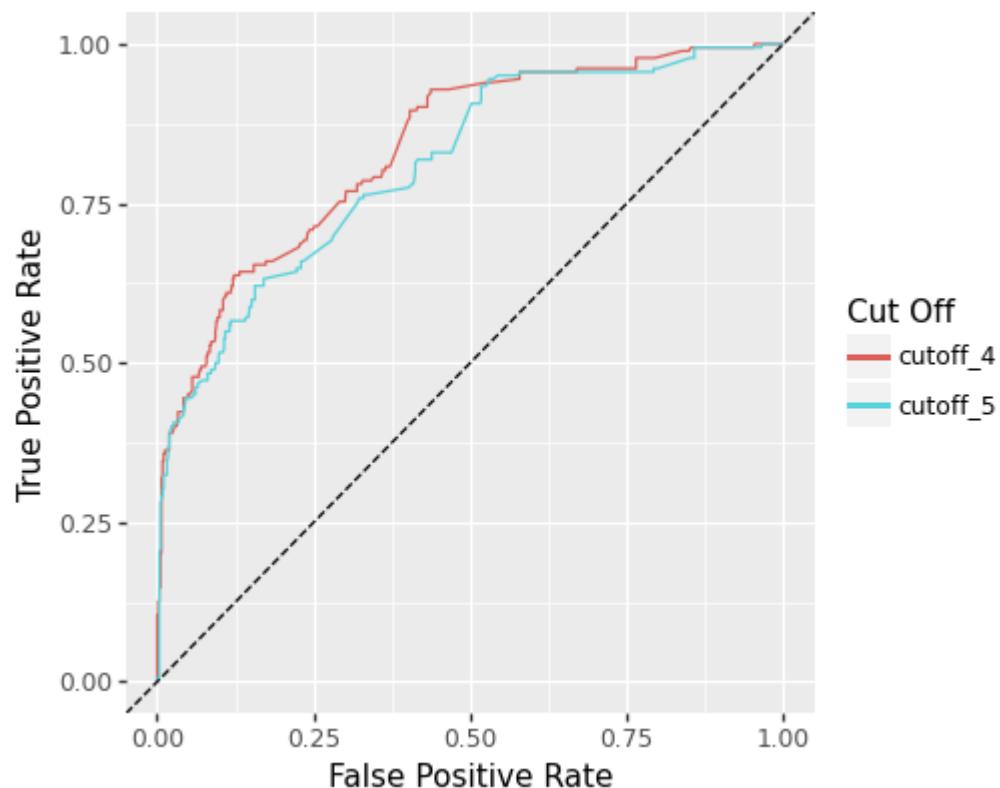
Summary Plots (`modeling_results/model_performance/`):

- SHAP beeswarm plots across all runs
- ROC curve comparisons across feature selection cutoffs
- Precision-Recall curve comparisons
- Hit rate and hit ratio curves

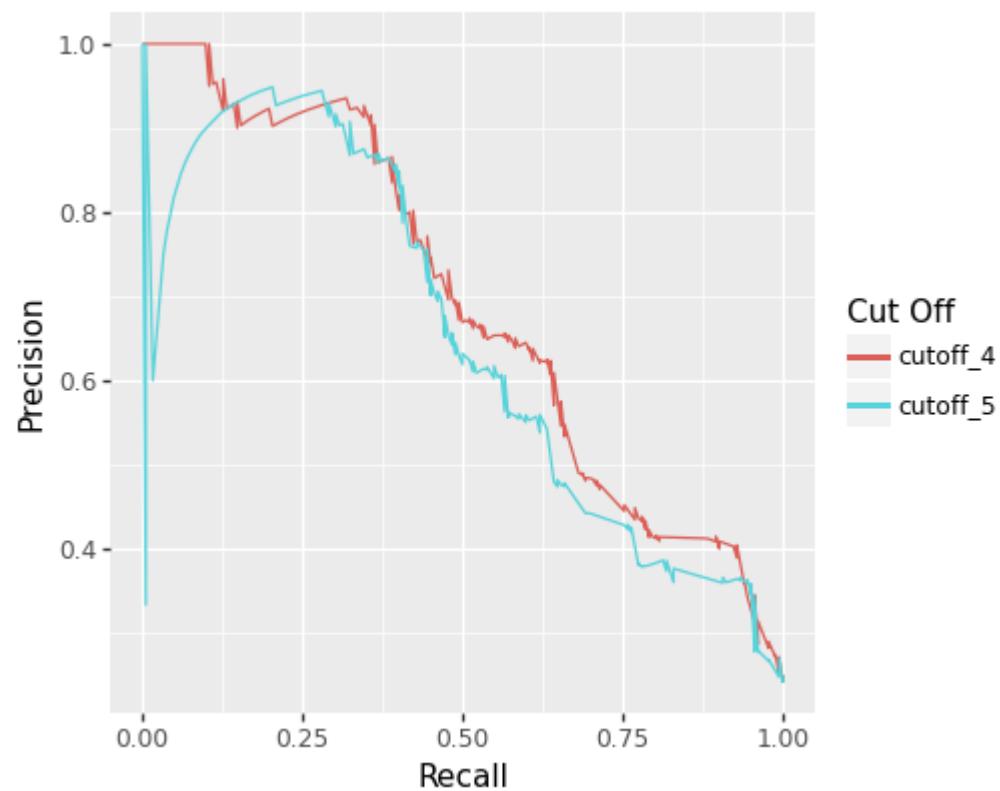
Example visualizations from the original README:



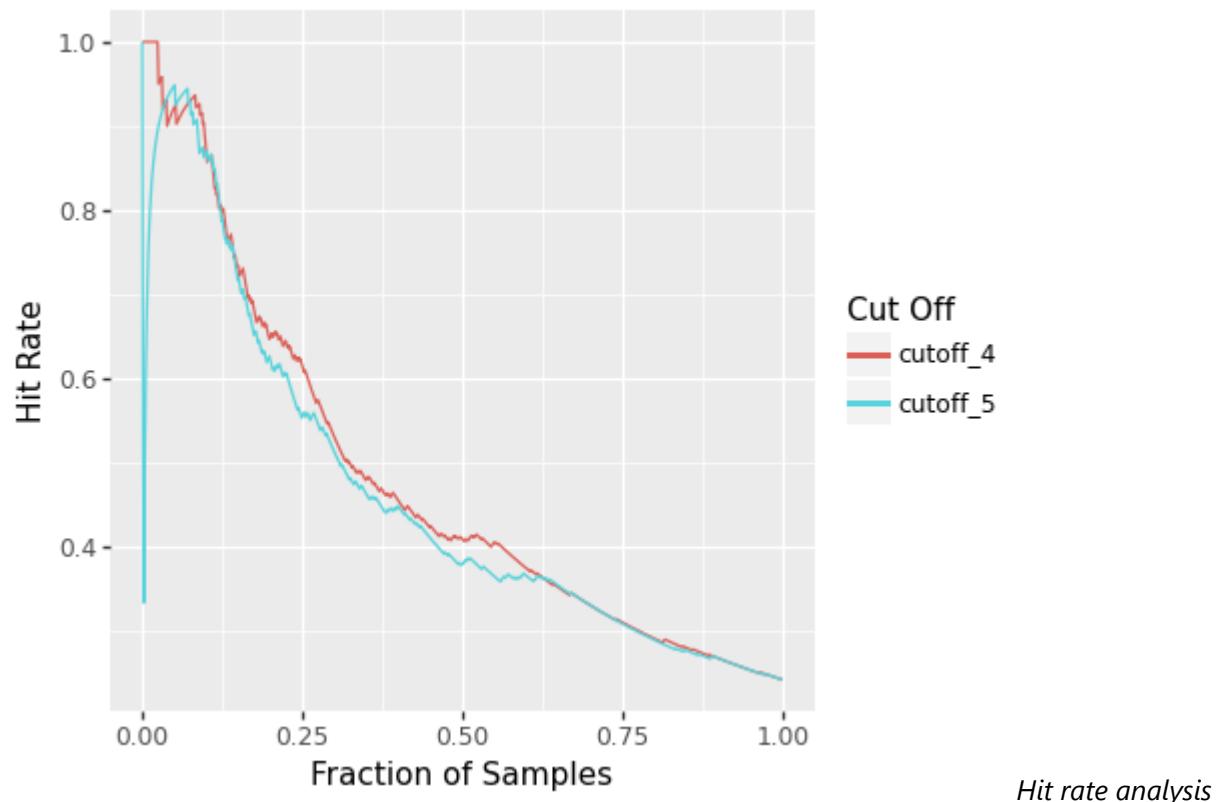
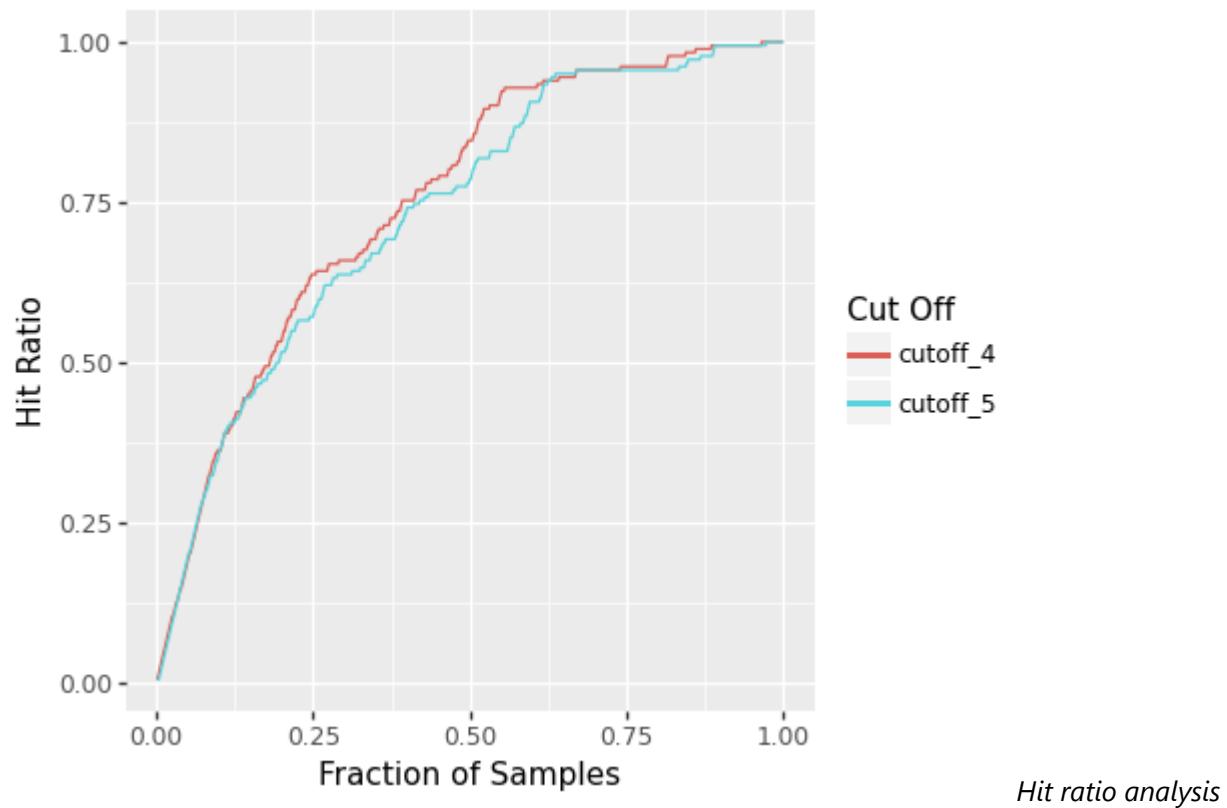
feature importance summary



ROC curves comparing different feature selection cutoffs



Precision-Recall curves

*Hit rate analysis**Hit ratio analysis*

Output Directory Structure

Protein Family Workflow

```
output_dir/
  └── clustering/
    ├── strain_db/          # MMseqs2 database
    └── strain_clusters.tsv # Cluster assignments
```

```

    └── phage_db/           # Phage database (optional)
        └── phage_clusters.tsv

    └── feature_tables/
        ├── strain_feature_table.csv
        ├── phage_feature_table.csv # If phages provided
        └── merged_feature_table.csv

        └── feature_map.csv

    └── feature_selection/
        ├── selected_features/
            ├── run_1_selected_features.csv
            └── ...
        └── feature_occurrence_counts.csv

        └── filtered_feature_tables/
            ├── cutoff_5_feature_table.csv
            └── cutoff_10_feature_table.csv
            └── ...

    └── modeling_results/
        ├── cutoff_5/
            ├── run_1/, run_2/, ...
            └── cutoff_5_combined_performance.csv

        ├── cutoff_10/
        └── model_performance/
            ├── shap_summary_plots/
            ├── roc_curve.png
            ├── pr_curve.png
            └── performance_comparison.csv

        └── models/
            ├── run_1/best_model.pkl
            └── ...

    └── predictions/
        ├── predicted_interactions.csv
        └── prediction_confidence.csv

    └── workflow_report.txt

```

K-mer Workflow

```

output_dir/
    └── clustering/           # Protein family clustering
    └── kmer_tables/
        ├── kmer_feature_table.csv
        ├── feature_map.csv
        ├── filtered_kmers.csv
        ├── feature2cluster.csv
        └── aa_sequences.faa

    └── feature_selection/
    └── modeling_results/
    └── kmer_workflow_report.txt

```

Python API

GenoPHI can also be used programmatically:

```
from genophi.workflows import (
    run_protein_family_workflow,
    run_kmer_workflow,
    run_modeling_workflow_from_feature_table,
    assign_predict_workflow
)

# Recommended: Protein family workflow
run_protein_family_workflow(
    input_path_strain="strain_fastas/",
    input_path_phage="phage_fastas/",
    phenotype_matrix="interactions.csv",
    output_dir="results/",
    threads=8,
    num_features=100,
    num_runs_fs=50,
    num_runs_modeling=100,
    method='rfe',
    filter_type='strain'
)

# _K_-mer workflow
run_kmer_workflow(
    input_strain_dir="strain_fastas/",
    input_phage_dir="phage_fastas/",
    phenotype_matrix="interactions.csv",
    output_dir="kmer_results/",
    k=5,
    threads=8,
    num_features=100
)

# Feature selection and modeling from any feature table
run_modeling_workflow_from_feature_table(
    full_feature_table="custom_features.csv",
    output_dir="modeling_results/",
    phenotype_column="your_phenotype",
    sample_column="your_sample_id",
    num_features=100,
    num_runs_fs=50,
    num_runs_modeling=100,
    method='rfe'
)

# Prediction workflow
assign_predict_workflow(
    input_dir="new_genomes/",
    mmseqs_db="results/tmp/strain/mmseqs_db",
    clusters_tsv="results/strain/clusters.tsv",
    feature_map="results/strain/features/selected_features.csv",
    model_dir="results/modeling_results/cutoff_*
```

```
    output_dir="predictions/",
    genome_type='strain',
    phage_feature_table_path="results/phage/features/feature_table.csv"
)
```

Advanced Usage

Custom Train-Test Splits

Control how data is split for model evaluation:

```
--filter_type none      # Random split
--filter_type strain    # Leave-strain-out (test on new strains)
--filter_type phage      # Leave-phage-out (test on new phages)
```

Recommendation: Use `--filter_type strain` for phage-host predictions to evaluate generalization to new strains.

Hyperparameter Tuning

Models use grid search for hyperparameter optimization. Default parameters are optimized for phage-host interaction prediction but can be customized in the code.

Handling Large Datasets

For large datasets, adjust memory and threading:

```
--max_ram 64 \          # Maximum RAM in GB
--threads 16 \            # Number of CPU threads
--clear_tmp               # Remove temporary files after completion
```

Regression Tasks

For continuous phenotypes:

```
--task_type regression \
--phenotype_column efficiency
```

GenoPHI will use appropriate regression metrics (RMSE, MAE, R²) instead of classification metrics.

Pre-processing Feature Clustering

Filter features by cluster presence before modeling:

```
--use_feature_clustering \
--feature_cluster_method hierarchical \
--feature_n_clusters 20 \
--feature_min_cluster_presence 2
```

This removes features that appear in fewer than `feature_min_cluster_presence` genome clusters.

Example Datasets

To help you get started, we recommend testing GenoPHI with:

Test Dataset Structure

Minimal test case with 10 strains, 5 phages:

```
test_data/
└── strains/
    ├── Strain_001.faa
    ├── Strain_002.faa
    └── ...
└── phages/
    ├── Phage_A.faa
    ├── Phage_B.faa
    └── ...
└── interactions.csv
```

Troubleshooting

Common Issues

Issue: MMseqs2 not found

```
Solution: Ensure MMseqs2 is installed and in your PATH
conda install -c bioconda mmseqs2
which mmseqs # Should show the path
```

Issue: Out of memory errors

```
Solution:
- Reduce --max_ram parameter
- Process fewer genomes at once
- Use --clear_tmp to remove intermediate files
- Increase system swap space
```

Issue: Duplicate protein IDs

Solution: GenoPHI automatically detects and prefixes duplicates with genome names
If you want to prevent this, ensure protein IDs are unique across all input files

Issue: No predictive features found

Solution:

- Try different feature selection methods (--method)
- Adjust num_features parameter
- Check that phenotype matrix has sufficient positive/negative examples
- Verify that interaction matrix matches genome filenames

Issue: Poor model performance

Solution:

- Increase num_runs_fs and num_runs_modeling for more robust results
- Try different feature selection methods
- Use --use_dynamic_weights for imbalanced features
- Enable --use_clustering for feature grouping
- Check data quality and ensure phenotype matrix is correct
- Try different clustering parameters (min_seq_id, coverage)

Issue: Models take too long to train

Solution:

- Reduce num_runs_modeling
- Reduce num_features
- Increase --threads parameter
- Use faster feature selection methods (select_k_best, chi_squared)

Frequently Asked Questions (FAQ)

General Questions

Q: Can GenoPHI be used for organisms other than phages and bacteria?

A: Yes! While designed for phage-host interactions, GenoPHI works with any protein sequences and phenotypes. It's been applied to bacteria, archaea, and eukaryotic microbes.

Q: How many genomes do I need for reliable predictions?

A: Minimum: ~20 strains and 20 phages with ~400 interactions. Recommended: 50+ strains, 50+ phages, 5000+ interactions for robust models.

Q: What file formats are required?

A: FASTA files (.faa) for protein sequences and CSV for phenotype matrices. See [Input Data Formats](#) for details.

Technical Questions

Q: What's the difference between protein family and k-mer approaches?

A: Protein families group similar full-length proteins (interpretable, captures protein-level patterns). K-mers analyze short amino acid sequences (high resolution, captures local patterns). The `full-workflow` combines both by extracting *k*-mers from predictive protein families.

Q: Should I use single-strain or phage-host mode?

A: Use phage-host mode for interaction prediction. Use single-strain mode for strain-level phenotypes (resistance, growth rate, etc.) where phage data isn't relevant.

Q: Which feature selection method should I use?

A: Start with RFE (balanced performance). Try SHAP for interpretability or SelectKBest for speed. Compare multiple methods for best results.

Q: How do I interpret SHAP plots?

A: SHAP beeswarm plots show feature importance. Features at the top are most important. Red dots = high feature values, blue = low. Position right of center = positive impact on prediction.

Q: Can I use custom features instead of protein families?

A: Yes! Use `select-and-train` with any feature table containing a phenotype column (metabolic pathways, gene presence/absence, etc.).

Q: How do I handle imbalanced datasets?

A: Use `--use_dynamic_weights` with `--weights_method inverse_frequency` to balance feature importance. CatBoost also has built-in class balancing.

Troubleshooting Questions

Q: Models perform poorly - what should I try?

A: (1) Increase num_runs for more robust estimates, (2) Try different clustering parameters, (3) Enable dynamic weighting, (4) Check data quality and phenotype matrix accuracy.

Q: How much RAM do I need?

A: Minimum 8 GB. Recommend 16+ GB for 50+ genomes, 32+ GB for 100+ genomes. Use `--max_ram` to limit memory usage.

Best Practices

1. **Start with recommended defaults** for initial analysis
2. **Use `--filter_type strain`** for phage-host predictions to evaluate generalization
3. **Run multiple iterations** (`num_runs_fs = 25, num_runs_modeling = 50`) for robust results
4. **Enable clustering** (`--use_clustering`) for correlated features
5. **Check data quality** before modeling - ensure phenotype matrix matches genome filenames
6. **Use SHAP plots** to understand which features drive predictions
7. **Compare multiple feature selection methods** to find optimal features
8. **For large datasets**, start with a subset to optimize parameters

Version History

v0.1.0 (Current)

- Initial release
- Protein family-based workflow with MMseqs2 clustering
- K-mer-based workflow with flexible k -mer lengths
- Multiple feature selection methods (RFE, SHAP, SelectKBest, Chi-squared, Lasso)
- CatBoost model training with hyperparameter optimization
- SHAP-based interpretability
- Support for classification and regression tasks
- Single-strain and phage-host prediction modes
- Unified CLI with 12 commands
- Comprehensive visualization outputs

Upcoming Features

- Web interface for prediction and visualization
- Docker container for easy deployment

Publication Datasets

The datasets used in the GenoPHI publication are included in the `data/` directory for reproducibility and benchmarking purposes.

```
data/
└── experimental_validation/
    ├── BASEL_ECOR_interaction_matrix.csv      # BASEL collection against ECOR
    strains for model validation
    └── ECOR27_TnSeq_high_fitness_genes.csv   # Filtered RB-TnSeq results
└── interaction_matrices/
    ├── ecoli_interaction_matrix.csv          # E. coli phage-host interactions
    ├── ecoli_interaction_matrix_subset.csv    # Smaller E. coli subset for testing
    ├── klebsiella1_interaction_matrix.csv     # Klebsiella dataset 1
    ├── klebsiella2_interaction_matrix.csv     # Klebsiella dataset 2
    ├── pseudomonas_interaction_matrix.csv     # Pseudomonas interactions
    └── vibrio_interaction_matrix.csv          # Vibrionaceae interactions
```

Citation

If you use GenoPHI in your research, please cite:

```
@article{noonan2025genophi,
  author = {Noonan, Avery J. C. and Moriniere, Lucas and Rivera-López, Edwin O. and Patel, Krish and Pena, Melina and Svab, Madeline and Kazakov, Alexey and Deutschbauer, Adam and Dudley, Edward G. and Matalik, Vivek K. and Arkin, Adam P.},
  title = {Phylogeny-agnostic strain-level prediction of phage-host interactions from genomes},
  year = {2025},
```

```
doi = {10.1101/2025.11.15.688630},  
publisher = {Cold Spring Harbor Laboratory},  
url = {https://www.biorxiv.org/content/10.1101/2025.11.15.688630v1},  
journal = {bioRxiv}  
}
```

Preprint: Noonan, A.J.C., Moriniere, L., Rivera-López, E.O., Patel, K., Pena, M., Svab, M., Kazakov, A., Deutschbauer, A., Dudley, E.G., Mutualik, V.K., & Arkin, A.P. (2025). Phylogeny-agnostic strain-level prediction of phage-host interactions from genomes. *bioRxiv*. <https://doi.org/10.1101/2025.11.15.688630>

Contributing

We welcome contributions to GenoPHI! Here's how you can help:

Ways to Contribute

- **Report bugs** or suggest features via [GitHub Issues](#)
- **Submit pull requests** for bug fixes or new features
- **Improve documentation** by fixing typos or adding examples
- **Share your use cases** and publications using GenoPHI
- **Test on different platforms** and report compatibility issues

Development Guidelines

1. Fork the repository and create a feature branch
2. Follow PEP 8 style guidelines for Python code
3. Add tests for new functionality
4. Update documentation as needed
5. Submit a pull request with a clear description of changes

Code of Conduct

We are committed to providing a welcoming and inclusive environment. Please be respectful and constructive in all interactions.

License

This software is available under the MIT License. See the [LICENSE](#) file for details.

This software is subject to Lawrence Berkeley National Laboratory copyright. The U.S. Government retains certain rights as this software was developed under funding from the U.S. Department of Energy.

Support

For questions, issues, or feature requests:

- Open an issue on [GitHub](#)
- Contact: Avery Noonan (averynoonan@gmail.com)

Acknowledgments

This was completed as part of the BRaVE Phage Foundry at Lawrence Berkeley National Laboratory which is supported by the U.S. Department of Energy, Office of Science, Office of Biological & Environmental Research under contract number DE-AC02-05CH11231. This work was also supported by the National Science Foundation (NSF) of the United States under grant award No. 2220735 (EDGE CMT: Predicting bacteriophage susceptibility from Escherichia coli genotype).